

## 5 Regularisation

### Contents

<b>5 Regularisation</b>	<b>1</b>
5.1 Addressing the issue of overfitting . . . . .	1
5.2 Cost function modification for regularisation . . . . .	2
5.3 Optimisation: Gradient descent for regularisation . . . . .	3
5.4 Optimisation: Normal equations for regularisation . . . . .	3
5.5 R Example . . . . .	4

#### Abstract

Regularisation is designed to address the problem of overfitting. **High bias** or **underfitting** is when the form of our hypothesis maps poorly to the trend of the data. It is usually caused by a function that is too simple or uses too few features (e.g. using a linear regression hypothesis when clearly the data appear quadratic).

At the other extreme, **overfitting** or **high variance** is caused by a hypothesis function that fits the available data but does not generalize well to predict new data. It is usually caused by a complicated function that creates a lot of unnecessary curves and angles unrelated to the data.

This terminology is applied to both linear and logistic regression.




Figure 5.1: Underfitting vs. "a good fit" vs. Overfitting

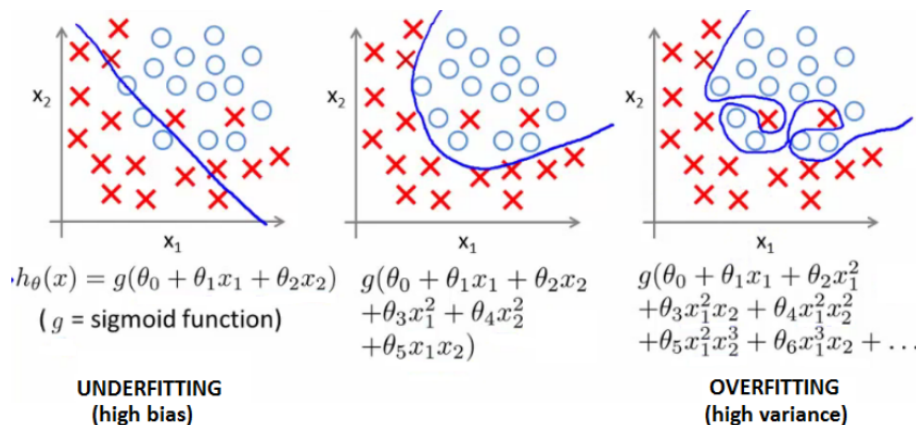


Figure 5.2: Logistic regression: Underfitting vs. "a good fit" vs. Overfitting

### 5.1 Addressing the issue of overfitting

There are 2 main options for addressing the issue of overfitting:

1. **Reduce the number of features.** This can be achieved by:
  - (a) *Manually* selecting which features to keep.

- (b) Using a model selection algorithm (studied later in the course).

Of course, reducing the number of features ultimately results in losing some information (we ideally select those features which minimize data loss).

## 2. Regularization

- (a) Keep all the features, but reduce the parameters  $\theta_j$ .  
 (b) This works well when we have a lot of features, each of which contributes a bit to predicting  $y$ .

Regularisation works well when we have a lot of slightly useful features.

## 5.2 Cost function modification for regularisation

To reduce the weights (or parameters) so that the following hypothesis function is more quadratic:

$$\theta_0 + \theta_1 x + \theta_2 x^2 + \theta_3 x^3 + \theta_4 x^4$$

we can eliminate the influence of  $\theta_3 x^3$  and  $\theta_4 x^4$ , without actually getting rid of these features or changing the form of our hypothesis. We can instead modify our cost function:

$$\min_{\theta} \frac{1}{2m} \sum_{i=1}^m \left( h_{\theta} \left( x^{(i)} \right) - y^{(i)} \right)^2 + 1000\theta_3^2 + 1000\theta_4^2$$




Figure 5.3: This modification of our cost function help to penalise  $\theta_3 x^3$  and  $\theta_4 x^4$ , so here we end up with  $\theta_3$  and  $\theta_4$  being close to zero (because the constants are massive) - so we're basically left with a quadratic function.

More generally, for regularisation we can modify the cost function:

**Definition 5.1.** The regularised objective function is defined as

$$\min_{\theta} \frac{1}{2m} \left[ \sum_{i=1}^m \left( h_{\theta} \left( x^{(i)} \right) - y^{(i)} \right)^2 + \lambda \sum_{j=1}^n \theta_j^2 \right]$$

By convention we don't penalize  $\theta_0$  - minimisation is from  $\theta_1$  (but in practice including  $\theta_0$  has negligible effect).

**Definition 5.2.**  $\lambda$  is the **regularisation parameter**, which determines how much the costs of our theta parameters are inflated.

*Remark.*  $\lambda$  controls a trade off between our two goals:

1. Fitting the training set well; we can smooth/make less curvy the output of our hypothesis function to reduce overfitting.
2. Keeping parameters small; WARNING: if  $\lambda \gg 1$  it may smooth out the function too much (i.e. effectively removing all terms) and cause underfitting.

We look at some automatic ways to select  $\lambda$  later in the course.

### 5.3 Optimisation: Gradient descent for regularisation

Because the  $\theta_0$  term is excluded from regularisation, the algorithm is modified in the following (again using calculus):

**Algorithm 1. Gradient Descent for regularisation**

repeat until convergence {

$$\theta_0 := \theta_0 - \frac{\alpha}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_0^{(i)}$$

$$\theta_j := \theta_j - \frac{\alpha}{m} \left[ \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)} + \lambda \theta_j \right], \text{ for } j = 1, 2, \dots, n$$

}

for  $\alpha$  some (positive) **learning rate** (aka. step size) and some initial guess for  $\theta$ .

*Remark.* Grouping the  $\theta_j$  terms, we see that

$$\theta_j = \theta_j \left( 1 - \alpha \frac{\lambda}{m} \right) - \frac{\alpha}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

$\frac{\alpha}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$  is exactly the same as the original gradient descent; the term  $\frac{\lambda \theta_j}{m}$  performs our regularisation. Noting that  $\alpha$  is small and  $m$  is large, we see that typically  $(1 - \alpha \frac{\lambda}{m}) < 1$ , in particular  $0.95 \leq (1 - \alpha \frac{\lambda}{m}) \leq 0.99$ . Thus, this in effect means  $\theta_j$  gets multiplied by 0.99 (reducing the value of  $\theta_j$  by some amount on every update).

### 5.4 Optimisation: Normal equations for regularisation

To add in regularisation to our normal equations, we use:

$$\theta = (X^T X + \lambda L)^{-1} X^T \mathbf{y}, \text{ where } L = \begin{bmatrix} 0 & & & \\ & 1 & & \\ & & \ddots & \\ & & & 1 \end{bmatrix} \in \mathbb{R}^{(n+1) \times (n+1)}$$

*Remark.* Intuitively,  $L$  is essentially the identity matrix multiplied by a single real number  $\lambda$  (not including  $x_0$ ).

## 5.5 R Example

```
> testidx <- which(1:nrow(Prestige)%%4==0)
> prestige_train <- Prestige[-testidx,]
> prestige_test <- Prestige[testidx,]

1 > library(glmnet)
2   > cv.fit <- cv.glmnet(as.matrix(prestige_train[,c(-4, -6)]), as.vector(prestige_train[,4]),
3   > plot(cv.fit)
4   > coef(cv.fit)
5   5 x 1 sparse Matrix of class "dgCMatrix"
6   1
7   (Intercept) 6.3876684930151
8   education   3.2111461944976
9   income      0.0009473793366
10  women       0.0000000000000
11  census      0.0000000000000
12  > prediction <- predict(cv.fit, newx=as.matrix(prestige_test[,c(-4, -6)]))
13  > cor(prediction, as.vector(prestige_test[,4]))
14  [,1]
15  1 0.9291181193
```

Figure 5.4: Regularisation in R. Note the second line reads `> cv.fit <- cv.glmnet(as.matrix(prestige_train[,c(-4, -6)]), as.vector(prestige_train[,4]), nlambda=100, alpha=0.7, family="gaussian")`

The `glmnet` library provides a cross-validation test to automatically choose the better  $\lambda$ :




Figure 5.5: Cross-validation plot for  $\lambda$ . It shows the best  $\lambda$  with minimal-root, mean-square error.